

Introduction to dplyr and other data manipulation techniques

Kiva Oken

April 5, 2017

The tidyverse

Group of 9 packages

dplyr, forcats, ggplot2, haven, purrr,
readr, stringr, tibble, tidyr

The tidyverse

Group of 9 packages

`dplyr`, `forcats`, `ggplot2`, `haven`, `purrr`,
`readr`, `stringr`, `tibble`, `tidyr`

The tidyverse

Group of 9 packages

`dplyr`, `forcats`, `ggplot2`, `haven`, `purrr`,
`readr`, `stringr`, `tibble`, `tidyr`

```
install.packages('tidyverse')
```

```
require(tidyverse)
```

Packages for another time

- forcats: Tools for dealing with factors
- ggplot2: Plotting
- haven: Read files from SAS, SPSS, Stata. NOT Matlab.
- purrr: Functional programming
- readr: Reading in rectangular data files
- stringr: Tools for dealing with strings.

- **Split** up a dataset
- **Apply** a function to each piece
- **Combine** pieces back together

Always returns a data frame!

Basic dplyr verbs

- **Filter:** filter rows, like `subset()`
- **Arrange:** reorder rows according to an index, like `df[order(index),]`
- **Select:** select certain columns, like `select` argument to `subset()` (also `$`, `df[,index]`)
- **Mutate:** Add new columns that are a function of other columns, like `transform()`
- **Summarize:** Collapses a data frame (or subset of a data frame) into a single row

filter() and slice()

```
olaf.assessments <- filter(assessment,  
                           recorder == 'JENSEN')
```


filter() and slice()

```
olaf.assessments <- filter(assessment,  
                           recorder == 'JENSEN')  
  
cod <- filter(stock, grepl('Gadus', scientificname))
```

filter() and slice()

```
olaf.assessments <- filter(assessment,  
                           recorder == 'JENSEN')  
  
cod <- filter(stock, grepl('Gadus', scientificname))  
  
slice(assessment, 5:8)
```

Basic dplyr verbs

- **Filter:** filter rows, like `subset()`
- **Arrange:** reorder rows according to an index, like `df[order(index),]`
- **Select:** select certain columns, like `select` argument to `subset()` (also `$`, `df[,index]`)
- **Mutate:** Add new columns that are a function of other columns, like `transform()`
- **Summarize:** Collapses a data frame (or subset of a data frame) into a single row

arrange()

```
arrange(olaf.assessments, daterecorded)
arrange(olaf.assessments, desc(daterecorded))
```

Basic dplyr verbs

- **Filter:** filter rows, like `subset()`
- **Arrange:** reorder rows according to an index, like `df[order(index),]`
- **Select:** select certain columns, like `select` argument to `subset()` (also `$`, `df[,index]`)
- **Mutate:** Add new columns that are a function of other columns, like `transform()`
- **Summarize:** Collapses a data frame (or subset of a data frame) into a single row

select() and rename()

```
select(olaf.assessments, stockid)
select(olaf.assessments, stock.id = stockid)
rename(olaf.assessments, stock.id = stockid)
select(olaf.assessments, assessid:stockid)
select(olaf.assessments, -recorder)
```

Basic dplyr verbs

- **Filter:** filter rows, like `subset()`
- **Arrange:** reorder rows according to an index, like `df[order(index),]`
- **Select:** select certain columns, like `select` argument to `subset()` (also `$`, `df[,index]`)
- **Mutate:** Add new columns that are a function of other columns, like `transform()`
- **Summarize:** Collapses a data frame (or subset of a data frame) into a single row

mutate() and transmute()

```
olaf.delay <- mutate(olaf.assessments,  
                      delay = dateloaded -  
                        daterecorded)  
select(olaf.delay, delay)  
  
transmute(olaf.assessments,  
           delay = dateloaded -  
             daterecorded)  
  
toothfish.ssb <- filter(timeseries, assessid ==  
                        'CCAMLR-ATOOTHFISHRS-1995-2007-JP',  
                        tsid == 'SSB-MT')  
mutate(toothfish.ssb,  
       zscore = (tsvalue - mean(tsvalue)) /  
                 sd(tsvalue))
```


Basic dplyr verbs

- **Filter:** filter rows, like `subset()`
- **Arrange:** reorder rows according to an index, like `df[order(index),]`
- **Select:** select certain columns, like `select` argument to `subset()` (also `$`, `df[,index]`)
- **Mutate:** Add new columns that are a function of other columns, like `transform()`
- **Summarize:** Collapses a data frame (or subset of a data frame) into a single row

summarize()

```
summarize(toothfish.ssb, mean(tsvalue, na.rm = TRUE))

##    mean(tsvalue, na.rm = TRUE)
## 1                67070.98

summarize(toothfish.ssb, n_distinct(tsvalue), n())

##    n_distinct(tsvalue) n()
## 1                11  13
```

summarize()

```
do_something <- function(vec) {  
  sum(vec, na.rm = TRUE)/5  
}  
summarize(toothfish.ssb, do_something(tsvalue))  
summarise(toothfish.ssb, do_something(tsvalue))
```

Aggregating functions

Accept vectors. Return scalars. (True?)

- Used for `summarize()`
- Base R: `min()`, `max()`, `mean()`, etc.
- From dplyr:
`n()`, `n_distinct()`, `first()`, `last()`, `nth()`
- Write your own, can use C++ for speed

And one more: Do()

- **Do:** general purpose verb to complement the specialized functions, like `dlply()`

```
assessors <- do(olaf.assessments,  
               assessor = unique(.$assessorid))  
assessors  
assessors$assessor  
class(assessors)  
  
distinct(olaf.assessments, assessorid)
```

Practice

- 1 Create a data frame in R that contains the time series data of Atlantic Amberjack using `filter()`. **Hint:** this stock is in `olaf.assessments`.
- 2 Using `filter()` and one other dplyr function, determine in which year Amberjack had the highest recruitment (R-E00).

Pipes! (%>%)

$$f(x) \%>\% g(y) \Leftrightarrow g(f(x), y)$$

Pipes! (%>%)

$$f(x) \%>\% g(y) \Leftrightarrow g(f(x), y)$$

Three ways to get the same result:

Pipes! (%>%)

$$f(x) \%>\% g(y) \Leftrightarrow g(f(x), y)$$

Three ways to get the same result:

```
1 x <- rnorm(100)
  x.mat <- matrix(x, nrow = 10)
  x.mns <- apply(x.mat, 1, mean)
```

Pipes! (%>%)

$$f(x) \%>\% g(y) \Leftrightarrow g(f(x), y)$$

Three ways to get the same result:

```
1 x <- rnorm(100)
  x.mat <- matrix(x, nrow = 10)
  x.mns <- apply(x.mat, 1, mean)
```

```
2 x.mns <- apply(matrix(rnorm(100), nrow = 10),
                     1, mean)
```

Pipes! (%>%)

$$f(x) \%>\% g(y) \Leftrightarrow g(f(x), y)$$

Three ways to get the same result:

```
1 x <- rnorm(100)
  x.mat <- matrix(x, nrow = 10)
  x.mns <- apply(x.mat, 1, mean)
```

```
2 x.mns <- apply(matrix(rnorm(100), nrow = 10),
                     1, mean)
```

```
3 x.mns <- rnorm(100) %>%
  matrix(nrow = 10) %>%
  apply(1, mean)
```

Pipes, continued

```
filter(timeseries,  
       stockid == 'GRAMBERSATLC',  
       tsid == 'SSB-MT') %>%  
  with(plot(tsyear, tsvalue, type = 'l',  
           xlab = 'Year', ylab = 'SSB (mt)',  
           main = stocklong[1]))  
  
lm(tsvalue ~ tsyear, data = toothfish.ssb) %>%  
  with(plot(fitted.values, residuals))
```

group_by()

The workhorse of dplyr, like `tapply()`, `ddply()`

```
toothfish <- filter(timeseries, assessid ==  
  'CCAMLR-ATOOTHFISHRS-1995-2007-JENSEN') %>%  
  select(tsid:tsvalue) %>%  
  group_by(tsid)  
  
summarize(toothfish, mn = mean(tsvalue, na.rm = TRUE),  
  stdev = sd(tsvalue, na.rm = TRUE))  
slice(toothfish, 1)  
mutate(toothfish, z.score =  
  (tsvalue - mean(tsvalue, na.rm=TRUE)) /  
  sd(tsvalue, na.rm=TRUE)) %>%  
  View()
```

Join

Join multiple data frames together based on a common variable
(e.g., species)

Join

Join multiple data frames together based on a common variable (e.g., species)

- **Inner join:** rows with matching values in both data frames, columns from both data frames
- **Left join:** all rows from first (left) data frame, columns from both data frames
- **Semi join:** rows with matching values in both data frames, columns from first data frame
- **Anti join:** rows from first data frame *without* matching values in second, columns from first data frame

Join example

```
select(stock, stockid, scientificname, commonname) %>%  
  inner_join(assessment) %>%  
  View()
```

```
## Joining, by = "stockid"
```


Tidy data

Untidy data (wide)

Control	Treatment
c_1	t_1
\vdots	\vdots
c_n	t_n

Tidy data (long)

Condition	Value
Control	c_1
\vdots	\vdots
Control	c_n
Treatment	t_1
\vdots	\vdots
Treatment	t_n

The tidyr package

Update of reshape2

- `gather()`: Wide to long data frame
- `spread()`: Long to wide data frame

```
wide.toothfish <- ungroup(toothfish) %>%  
  mutate(tsid = gsub('-', '_', tsid)) %>%  
  spread(key = 'tsid', value = 'tsvalue')  
long.toothfish <- gather(wide.toothfish,  
                        key = tsid, value = tsvalue,  
                        BdivBmgttouse_dimensionless:  
                        Utouse_index)
```

Practice II

- 1 Create a data frame in R of data for Pacific herring (*Clupea pallasii*) that is grouped by stock and population metric (SSB, recruitment, etc.). You will need to join information from all three of the data tables in the RAM database to do this.
- 2 Using your data frame, calculate the mean and standard deviation of each population metric (SSB, recruitment, etc.) for each area. Note that the database contains NAs.
- 3 Plot the time series of spawning stock biomass of Pacific herring to compare across regions using either `do()` with base graphics or `ggplot()`.
- 4 Bonus: Color the lines produced above by exploitation rate (ER-ratio). You may want to use `tidyr`.

The end

Further resources:

- Package vignettes for `dplyr`, `tidyr`
- www.tidyverse.org
- Trevor Branch Super Advanced R course webpage
- Google, Stack Exchange, etc.